

CSP554—Big Data Technologies

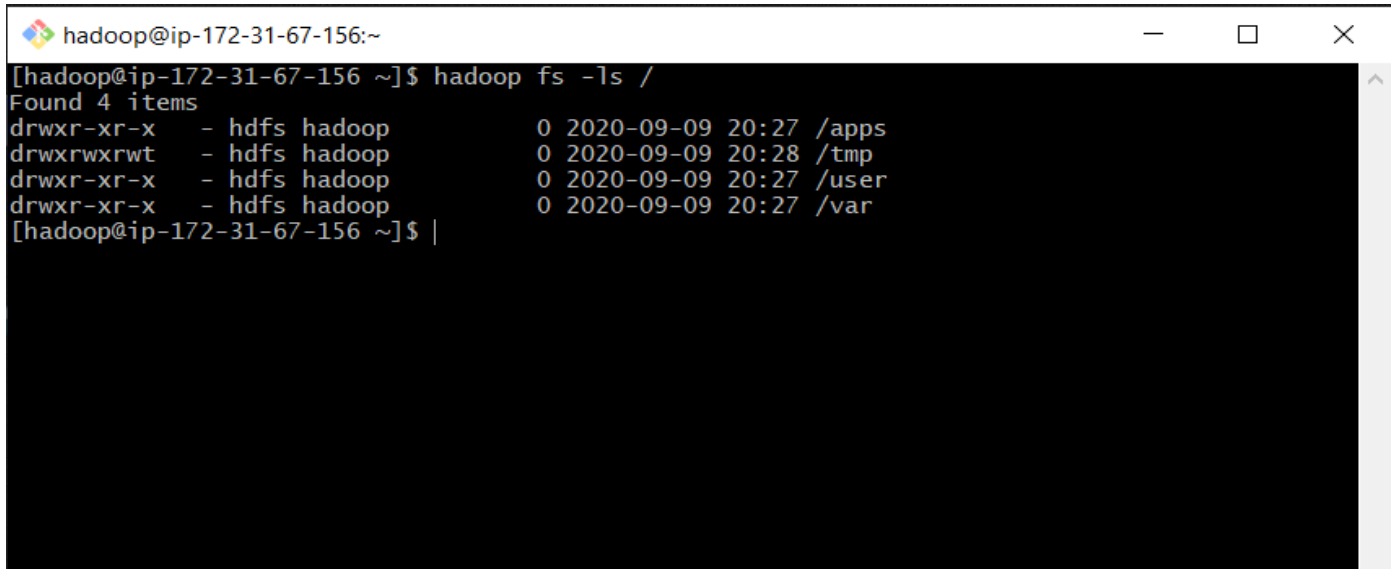
Assignment #2 (Modules 02a & 02b, 20 points)

9. (2 points) Execute the following hdfs command to list the files or directories that are listed (also indicating which is a file and which a directory):

hadoop fs -ls /

Ans:

Command: `hadoop fs -ls /`

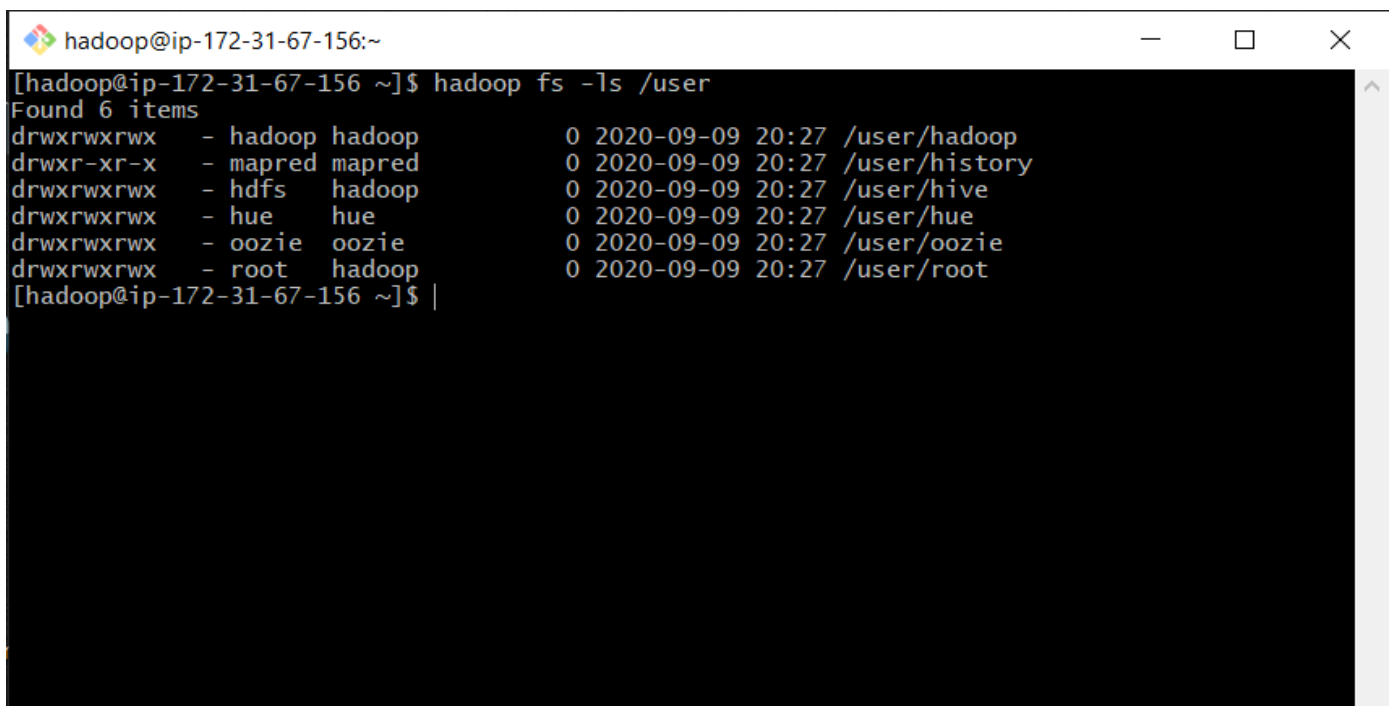


```
hadoop@ip-172-31-67-156:~  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /  
Found 4 items  
drwxr-xr-x - hdfs hadoop 0 2020-09-09 20:27 /apps  
drwxrwxrwt - hdfs hadoop 0 2020-09-09 20:28 /tmp  
drwxr-xr-x - hdfs hadoop 0 2020-09-09 20:27 /user  
drwxr-xr-x - hdfs hadoop 0 2020-09-09 20:27 /var  
[hadoop@ip-172-31-67-156 ~]$ |
```

10. (2 points) Execute a command (you needed to figure out which one) to list the files and directories under the hdfs directory listed below: /user

Ans:

Command: `hadoop fs -ls /user`



```
hadoop@ip-172-31-67-156:~  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user  
Found 6 items  
drwxrwxrwx - hadoop hadoop 0 2020-09-09 20:27 /user/hadoop  
drwxr-xr-x - mapred mapred 0 2020-09-09 20:27 /user/history  
drwxrwxrwx - hdfs hadoop 0 2020-09-09 20:27 /user/hive  
drwxrwxrwx - hue hue 0 2020-09-09 20:27 /user/hue  
drwxrwxrwx - oozie oozie 0 2020-09-09 20:27 /user/oozie  
drwxrwxrwx - root hadoop 0 2020-09-09 20:27 /user/root  
[hadoop@ip-172-31-67-156 ~]$ |
```

11. (2 points) Execute a command to create the following directory: /user/csp554

Ans:

Command: `hadoop fs -mkdir /user/csp554`

```
hadoop@ip-172-31-67-156:~  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user  
Found 6 items  
drwxrwxrwx - hadoop hadoop      0 2020-09-09 20:27 /user/hadoop  
drwxr-xr-x - mapred mapred      0 2020-09-09 20:27 /user/history  
drwxrwxrwx - hdfs hadoop      0 2020-09-09 20:27 /user/hive  
drwxrwxrwx - hue hue          0 2020-09-09 20:27 /user/hue  
drwxrwxrwx - oozie oozie      0 2020-09-09 20:27 /user/oozie  
drwxrwxrwx - root hadoop      0 2020-09-09 20:27 /user/root  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -mkdir /user/csp554  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user  
Found 7 items  
drwxr-xr-x - hadoop hadoop      0 2020-09-09 21:23 /user/csp554  
drwxrwxrwx - hadoop hadoop      0 2020-09-09 20:27 /user/hadoop  
drwxr-xr-x - mapred mapred      0 2020-09-09 20:27 /user/history  
drwxrwxrwx - hdfs hadoop      0 2020-09-09 20:27 /user/hive  
drwxrwxrwx - hue hue          0 2020-09-09 20:27 /user/hue  
drwxrwxrwx - oozie oozie      0 2020-09-09 20:27 /user/oozie  
drwxrwxrwx - root hadoop      0 2020-09-09 20:27 /user/root  
[hadoop@ip-172-31-67-156 ~]$ |
```

12. (2 points) Execute a command to create the following directory: /user/csp554-2

Ans:

Command: `hadoop fs -mkdir /user/csp554-2`

```
hadoop@ip-172-31-67-156:~  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user  
Found 7 items  
drwxr-xr-x - hadoop hadoop      0 2020-09-09 21:23 /user/csp554  
drwxrwxrwx - hadoop hadoop      0 2020-09-09 20:27 /user/hadoop  
drwxr-xr-x - mapred mapred      0 2020-09-09 20:27 /user/history  
drwxrwxrwx - hdfs hadoop      0 2020-09-09 20:27 /user/hive  
drwxrwxrwx - hue hue          0 2020-09-09 20:27 /user/hue  
drwxrwxrwx - oozie oozie      0 2020-09-09 20:27 /user/oozie  
drwxrwxrwx - root hadoop      0 2020-09-09 20:27 /user/root  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -mkdir /user/csp554-2  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user  
Found 8 items  
drwxr-xr-x - hadoop hadoop      0 2020-09-09 21:23 /user/csp554  
drwxr-xr-x - hadoop hadoop      0 2020-09-09 21:24 /user/csp554-2  
drwxrwxrwx - hadoop hadoop      0 2020-09-09 20:27 /user/hadoop  
drwxr-xr-x - mapred mapred      0 2020-09-09 20:27 /user/history  
drwxrwxrwx - hdfs hadoop      0 2020-09-09 20:27 /user/hive  
drwxrwxrwx - hue hue          0 2020-09-09 20:27 /user/hue  
drwxrwxrwx - oozie oozie      0 2020-09-09 20:27 /user/oozie  
drwxrwxrwx - root hadoop      0 2020-09-09 20:27 /user/root  
[hadoop@ip-172-31-67-156 ~]$ |
```

13. (2 points) Execute a command that copies a given local file to the given hdfs directory :

Source local file: /home/hadoop/myname.txt

Destination HDFS directory: /user/csp554

Ans:

Command: `hadoop fs -copyFromLocal /home/hadoop/akashtanwani.txt /user/csp554`

```
hadoop@ip-172-31-67-156:~  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user  
Found 8 items  
drwxr-xr-x - hadoop hadoop 0 2020-09-09 21:35 /user/csp554  
drwxr-xr-x - hadoop hadoop 0 2020-09-09 21:24 /user/csp554-2  
drwxrwxrwx - hadoop hadoop 0 2020-09-09 20:27 /user/hadoop  
drwxr-xr-x - mapred mapred 0 2020-09-09 20:27 /user/history  
drwxrwxrwx - hdfs hadoop 0 2020-09-09 20:27 /user/hive  
drwxrwxrwx - hue hue 0 2020-09-09 20:27 /user/hue  
drwxrwxrwx - oozie oozie 0 2020-09-09 20:27 /user/oozie  
drwxrwxrwx - root hadoop 0 2020-09-09 20:27 /user/root  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -copyFromLocal /home/hadoop/akashtanwani.txt /user/csp554  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user/csp554  
Found 1 items  
-rw-r--r-- 1 hadoop hadoop 23 2020-09-09 21:36 /user/csp554/akashtanwani.txt  
[hadoop@ip-172-31-67-156 ~]$ |
```

14. (2 points) Copy a file from one hdfs directory to another hdfs directory and write down the command

Source hdfs file: /user/csp554/myname.txt

Destination HDFS directory: /user/csp554-2

Ans:

Command: `hadoop fs -cp /user/csp554/akashtanwani.txt /user/csp554-2`

```
hadoop@ip-172-31-67-156:~  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -cp /user/csp554/akashtanwani.txt /user/csp554-2  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user/csp554-2  
Found 1 items  
-rw-r--r-- 1 hadoop hadoop 23 2020-09-09 21:43 /user/csp554-2/akashtanwani.txt  
[hadoop@ip-172-31-67-156 ~]$ |
```

15. (2 points) Copy the object myid.txt you uploaded to an S3 bucket into the Hadoop master node Linux file system.

The actual object includes your student id as above.

Ans:

Command: `aws s3 cp s3://a20448831-mybucket/A20448831.txt /home/Hadoop`
`ls /home/hadoop`

```
hadoop@ip-172-31-67-156:~  
[hadoop@ip-172-31-67-156 ~]$ aws s3 cp s3://a20448831-mybucket/A20448831.txt /home/hadoop  
download: s3://a20448831-mybucket/A20448831.txt to ./A20448831.txt  
[hadoop@ip-172-31-67-156 ~]$ ls /home/hadoop  
A20448831.txt akashtanwani.txt  
[hadoop@ip-172-31-67-156 ~]$ |
```

16. (2 points) Copy the same object myid.txt you created in an S3 bucket into HDFS into the directory /users/csp554

```
hadoop fs -cp s3://mybucket/myid.txt hdfs:///user/csp554-2
```

Note, the three slashes after the “hdfs:”

Ans:

Command: `hadoop fs -cp s3://a20448831-mybucket/A20448831.txt hdfs:///user/csp554-2`

`hadoop fs -ls /user/csp554-2`

```
hadoop@ip-172-31-67-156:~  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -cp s3://a20448831-mybucket/A20448831.txt hdfs:///user/csp554-2  
20/09/10 00:03:26 INFO s3n.S3NativeFileSystem: Opening 's3://a20448831-mybucket/A20448831.txt' for reading  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user/csp554-2  
Found 2 items  
-rw-r--r-- 1 hadoop hadoop 21 2020-09-10 00:03 /user/csp554-2/A20448831.txt  
-rw-r--r-- 1 hadoop hadoop 23 2020-09-09 21:43 /user/csp554-2/akasthanwani.txt  
[hadoop@ip-172-31-67-156 ~]$ |
```

17. (2 points) Execute a command to show the contents of the myid.txt file in the hdfs directory /user/csp554-2

Clue: look up about how to use the “cat” command in the file system shell document.

Ans:

Command: `hadoop fs -cat /user/csp554-2/A20448831.txt`

```
hadoop@ip-172-31-67-156:~  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -cat /user/csp554-2/A20448831.txt  
this is the id file  
[hadoop@ip-172-31-67-156 ~]$ |
```

18. (2 points) Execute a command to remove the myid.txt file in the hdfs directory /user/csp554-2

Ans:

Command: `hadoop fs -rm /user/csp554-2/A20448831.txt`

`hadoop fs -ls /user/csp554-2`

```
MINGW64;C:/Users/Akash Tanwani  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user/csp554-2  
Found 2 items  
-rw-r--r-- 1 hadoop hadoop 21 2020-09-10 00:03 /user/csp554-2/A20448831.txt  
-rw-r--r-- 1 hadoop hadoop 23 2020-09-09 21:43 /user/csp554-2/akasthanwani.txt  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -rm /user/csp554-2/A20448831.txt  
Deleted /user/csp554-2/A20448831.txt  
[hadoop@ip-172-31-67-156 ~]$ hadoop fs -ls /user/csp554-2  
Found 1 items  
-rw-r--r-- 1 hadoop hadoop 23 2020-09-09 21:43 /user/csp554-2/akasthanwani.txt  
[hadoop@ip-172-31-67-156 ~]$ exit  
logout  
Connection to ec2-3-237-22-252.compute-1.amazonaws.com closed.  
Akash@Akash MINGW64 ~  
$ |
```